

Tasty Search!

Build a search engine to search gourmet food reviews data and return the top K reviews that have the highest overlap with the input query. Requirements are as follows:

- Use the dataset available at <http://snap.stanford.edu/data/web-FineFoods.html>. Use [this alternate download link](#) in case of issues.
- Randomly sample 100K reviews from the above set for the index. In case of memory issues, downsample to 50K.
- Build an api that returns the top K highest scoring documents for any given query, where the score is defined by the following:

$$Score(D, Q) = \frac{|Q \cap D|}{|Q|}$$

(i.e. # tokens matching between Query(Q) & Document(D) normalized by query length -- the number of tokens in the given query).

e.g. given $Q = \{ \text{cat, processed, bad, good} \}$ and a document described by the following:

```
product/productId: B001E4KFG0
review/userId: A3SGXH7AUHU8GW
review/profileName: delmartian
review/helpfulness: 1/1
review/score: 5.0
review/time: 1303862400
review/summary: Good Quality Dog Food
review/text: I have bought several of the Vitality canned dog food
products and have found them all to be of good quality. The
product looks more like a stew than a processed meat and it smells
better. My Labrador is finicky and she appreciates this product
better than most.
```

$$\text{Score} = 2 / 4 = 0.5$$

- Queries are input as a set of tokens (words) and multiple occurrences of a given token in the text are counted as a single hit. You can assume queries are between 1 - 10 tokens. Ties between scores can be resolved using review/score for the document.
- Set K to 20
- Expose a simple REST interface for search. A simple web form to perform search on the server itself would be great.

Load test and performance

Using tokens from the data set, generate a random list of 100K queries varying length and the tokens. Using this random queryset, loadtest your server and report baseline throughput and latency numbers.

Now, improve response latency by 20% over the above baseline using the same loadtest setup. Write a short explanation of your approach(es) and the improvement achieved.

Deliverables

- Link to Github repository with the code along with instructions on running it. Please also include the generated queryset.
- Bonus points for including a running instance on AWS!
- Write up explaining the performance improvement and the numbers achieved.
- Please feel free to use any language you are comfortable with.